

# DATA 622 Assignment 1

CUNY: Spring 2021

Philip Tanofsky

14 February 2021

```
library(palmerpenguins)
library(dplyr)
library(ggplot2)
library(tidyr)
library(caret)
library(MASS)
library(pROC)
library(nnet) # Used for multinomial logistic regression
library(stargazer)
theme_set(theme_minimal())

ds <- penguins

head(ds)
```

```
## # A tibble: 6 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex
##   <fct>   <fct>         <dbl>         <dbl>         <int>         <int> <fct>
## 1 Adelie  Torge~           39.1           18.7           181           3750 male
## 2 Adelie  Torge~           39.5           17.4           186           3800 fema~
## 3 Adelie  Torge~           40.3            18           195           3250 fema~
## 4 Adelie  Torge~           NA            NA            NA            NA <NA>
## 5 Adelie  Torge~           36.7           19.3           193           3450 fema~
## 6 Adelie  Torge~           39.3           20.6           190           3650 male
## # ... with 1 more variable: year <int>
```

```
summary(ds)
```

```
##      species      island  bill_length_mm  bill_depth_mm
## Adelie   :152  Biscoe   :168  Min.   :32.10  Min.   :13.10
## Chinstrap: 68  Dream    :124  1st Qu.:39.23  1st Qu.:15.60
## Gentoo   :124  Torgersen: 52  Median :44.45  Median :17.30
##                                     Mean   :43.92  Mean   :17.15
##                                     3rd Qu.:48.50  3rd Qu.:18.70
##                                     Max.   :59.60  Max.   :21.50
##                                     NA's   :2      NA's   :2
## flipper_length_mm  body_mass_g      sex      year
## Min.   :172.0      Min.   :2700  female:165  Min.   :2007
## 1st Qu.:190.0      1st Qu.:3550  male  :168  1st Qu.:2007
```

```
## Median :197.0      Median :4050      NA's   : 11      Median :2008
## Mean   :200.9      Mean   :4202                      Mean   :2008
## 3rd Qu.:213.0      3rd Qu.:4750                      3rd Qu.:2009
## Max.   :231.0      Max.   :6300                      Max.   :2009
## NA's   :2          NA's   :2
```


```
dim(ds)
```

```
## [1] 344    8
```

```
glimpse(ds)
```

```
## Rows: 344
## Columns: 8
## $ species      <fct> Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, A...
## $ island       <fct> Torgersen, Torgersen, Torgersen, Torgersen, Torge...
## $ bill_length_mm <dbl> 39.1, 39.5, 40.3, NA, 36.7, 39.3, 38.9, 39.2, 34....
## $ bill_depth_mm <dbl> 18.7, 17.4, 18.0, NA, 19.3, 20.6, 17.8, 19.6, 18....
## $ flipper_length_mm <int> 181, 186, 195, NA, 193, 190, 181, 195, 193, 190, ...
## $ body_mass_g   <int> 3750, 3800, 3250, NA, 3450, 3650, 3625, 4675, 347...
## $ sex          <fct> male, female, female, NA, female, male, female, m...
## $ year         <int> 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2...
```

```
visdat::vis_dat(ds)
```



```
# Penguins data has three factor variables
```

```
ds %>%  
  dplyr::select(where(is.factor)) %>%  
  glimpse()
```

```
## Rows: 344
```

```
## Columns: 3
```

```
## $ species <fct> Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, Ade...
```

```
## $ island <fct> Torgersen, Torgersen, Torgersen, Torgersen, Torgersen, Torg...
```

```
## $ sex <fct> male, female, female, NA, female, male, female, male, NA, N...
```

```
# Count penguins for each species / island
```

```
ds %>%  
  count(species, island, .drop=F)
```

```
## # A tibble: 9 x 3
```

```
##   species    island      n
```

```
##   <fct>      <fct>    <int>
```

```
## 1 Adelie    Biscoe      44
```

```
## 2 Adelie    Dream       56
```

```
## 3 Adelie    Torgersen    52
```

```
## 4 Chinstrap Biscoe       0
```

```
## 5 Chinstrap Dream       68
```


```
## 6 Chinstrap Torgersen    0
```

```
## 7 Gentoo    Biscoe     124
```

```
## 8 Gentoo    Dream       0
```

```
## 9 Gentoo    Torgersen    0
```


```
ggplot(ds, aes(x = island, fill = species)) +  
  geom_bar(alpha = 0.8) +  
  scale_fill_manual(values = c("darkorange", "purple", "cyan4"),  
                    guide = F) +  
  theme_minimal() +  
  facet_wrap(~species, ncol = 1) +  
  coord_flip()
```



```
# Count penguins for each species / sex
ds %>%
  count(species, sex, .drop = F)
```

```
## # A tibble: 8 x 3
##   species sex      n
##   <fct>   <fct> <int>
## 1 Adelie female   73
## 2 Adelie male    73
## 3 Adelie <NA>      6
## 4 Chinstrap female  34
## 5 Chinstrap male   34
## 6 Gentoo female   58
## 7 Gentoo male    61
## 8 Gentoo <NA>      5
```


```
ggplot(ds, aes(x = sex, fill = species)) +
  geom_bar(alpha = 0.8) +
  scale_fill_manual(values = c("darkorange", "purple", "cyan4"),
                    guide = F) +
  theme_minimal() +
  facet_wrap(~species, ncol = 1) +
  coord_flip()
```



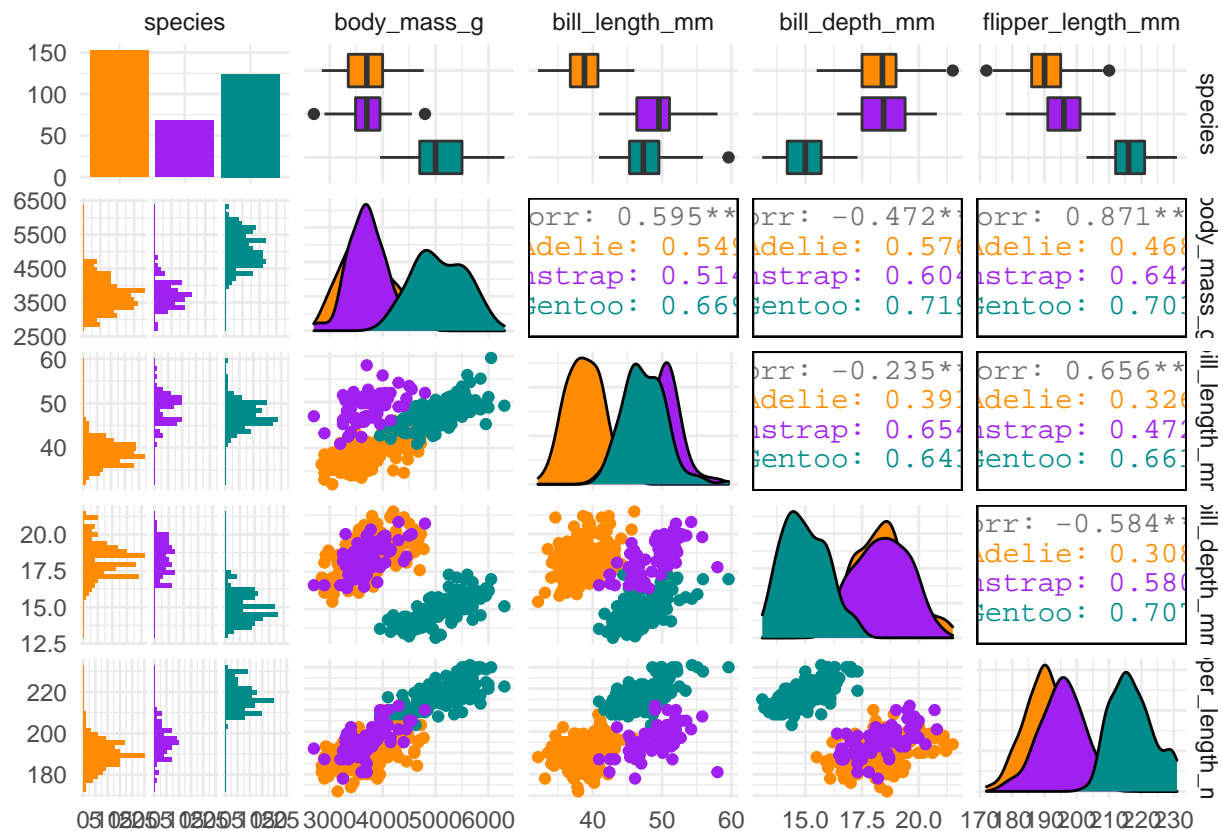
```
# Penguins data also has four continuous variables, making six unique scatterplots
ds %>%
  dplyr::select(body_mass_g, ends_with("_mm")) %>%
  glimpse()
```

```
## Rows: 344
## Columns: 4
## $ body_mass_g      <int> 3750, 3800, 3250, NA, 3450, 3650, 3625, 4675, 347...
## $ bill_length_mm   <dbl> 39.1, 39.5, 40.3, NA, 36.7, 39.3, 38.9, 39.2, 34....
## $ bill_depth_mm    <dbl> 18.7, 17.4, 18.0, NA, 19.3, 20.6, 17.8, 19.6, 18....
## $ flipper_length_mm <int> 181, 186, 195, NA, 193, 190, 181, 195, 193, 190, ...
```

```
# Scatterplot example 1: penguin flipper length versus body mass
ggplot(data = penguins, aes(x = flipper_length_mm, y = body_mass_g)) +
  geom_point(aes(color = species,
                 shape = species),
            size = 2) +
  scale_color_manual(values = c("darkorange", "darkorchid", "cyan4"))
```



```
ds %>%
  dplyr::select(species, body_mass_g, ends_with("_mm")) %>%
  GGally::ggpairs(aes(color = species)) +
  scale_color_manual(values = c("darkorange", "purple", "cyan4")) +
  scale_fill_manual(values = c("darkorange", "purple", "cyan4"))
```



```
# Create dataset for binary logistic regression: species Gentoo or Not
data_binary <- penguins
```

```
# Only use complete instances ... actually come back to this as I don't want to exclude because of sex
train_data_binary <- na.omit(data_binary)
```

```
dim(train_data_binary)
```

```
## [1] 333 8
```

Based on the result, 11 rows are removed, which would equal the number of NAs in variable sex

## Binary Logistic Regression

```
# Create new column
```

```
train_data_binary$gentoo <- ifelse(train_data_binary$species=="Gentoo", 1, 0)
```

```
summary(train_data_binary)
```

```
##      species      island  bill_length_mm  bill_depth_mm
## Adelie   :146  Biscoe   :163    Min.   :32.10    Min.   :13.10
## Chinstrap: 68   Dream   :123    1st Qu.:39.50    1st Qu.:15.60
```

```
## Gentoo :119 Torgersen: 47 Median :44.50 Median :17.30
## Mean :43.99 Mean :17.16
## 3rd Qu.:48.60 3rd Qu.:18.70
## Max. :59.60 Max. :21.50
## flipper_length_mm body_mass_g sex year gentoo
## Min. :172 Min. :2700 female:165 Min. :2007 Min. :0.0000
## 1st Qu.:190 1st Qu.:3550 male :168 1st Qu.:2007 1st Qu.:0.0000
## Median :197 Median :4050 Median :2008 Median :0.0000
## Mean :201 Mean :4207 Mean :2008 Mean :0.3574
## 3rd Qu.:213 3rd Qu.:4775 3rd Qu.:2009 3rd Qu.:1.0000
## Max. :231 Max. :6300 Max. :2009 Max. :1.0000
```

```
# Drop species column, as now just using gentoo column as Y variable
```

```
drops <- c("species")
train_data_binary <- train_data_binary[ , !(names(train_data_binary) %in% drops)]
summary(train_data_binary)
```

```
## island bill_length_mm bill_depth_mm flipper_length_mm
## Biscoe :163 Min. :32.10 Min. :13.10 Min. :172
## Dream :123 1st Qu.:39.50 1st Qu.:15.60 1st Qu.:190
## Torgersen: 47 Median :44.50 Median :17.30 Median :197
## Mean :43.99 Mean :17.16 Mean :201
## 3rd Qu.:48.60 3rd Qu.:18.70 3rd Qu.:213
## Max. :59.60 Max. :21.50 Max. :231
## body_mass_g sex year gentoo
## Min. :2700 female:165 Min. :2007 Min. :0.0000
## 1st Qu.:3550 male :168 1st Qu.:2007 1st Qu.:0.0000
## Median :4050 Median :2008 Median :0.0000
## Mean :4207 Mean :2008 Mean :0.3574
## 3rd Qu.:4775 3rd Qu.:2009 3rd Qu.:1.0000
## Max. :6300 Max. :2009 Max. :1.0000
```

```
set.seed(123)
trainIndex <- createDataPartition(train_data_binary$gentoo, p = 0.7, list = FALSE, times = 1)
train <- train_data_binary[trainIndex,]
test <- train_data_binary[-trainIndex,]
```

```
model1 <- glm(gentoo ~ ., data = train, family = "binomial"(link="logit")) %>% stepAIC(trace=F, direction="both")
summary(model1)
```

```
##
## Call:
## glm(formula = gentoo ~ bill_depth_mm + flipper_length_mm, family = binomial(link = "logit"),
## data = train)
##
## Deviance Residuals:
## Min 1Q Median 3Q Max
## -6.826e-05 -2.100e-08 -2.100e-08 2.100e-08 6.510e-05
##
## Coefficients:
```



```
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -428.130  474946.766  -0.001    0.999
## bill_depth_mm   -14.834  12021.243  -0.001    0.999
## flipper_length_mm  3.274   1957.819   0.002    0.999
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 3.0884e+02  on 233  degrees of freedom
## Residual deviance: 9.5967e-09  on 231  degrees of freedom
## AIC: 6
##
## Number of Fisher Scoring iterations: 25
```

```
#mu<-predict(model1, type = "response")

# calculate AIC
mod1AIC <- model1$aic

## use the test data set to make predicts and calculate metrics from the confusion matrix
mod1.predict.probs <- predict.glm(model1, type="response", newdata=test)

mod1.predict.manual <- ifelse(mod1.predict.probs > 0.5, '1','0')
attach(test)

mod1.predict.manual
```

```
##   1   2   3   4   5   6   7   8   9  10  11  12  13  14  15  16  17  18  19  20
## "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
##  21  22  23  24  25  26  27  28  29  30  31  32  33  34  35  36  37  38  39  40
## "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
##  41  42  43  44  45  46  47  48  49  50  51  52  53  54  55  56  57  58  59  60
## "0" "0" "0" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1"
##  61  62  63  64  65  66  67  68  69  70  71  72  73  74  75  76  77  78  79  80
## "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "1" "0" "0" "0" "0" "0"
##  81  82  83  84  85  86  87  88  89  90  91  92  93  94  95  96  97  98  99
## "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0" "0"
```

```
test$gentoo
```

```
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## [39] 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0
## [77] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

```
# now can use the caret function
cm.var <- caret::confusionMatrix(factor(mod1.predict.manual), factor(test$gentoo), positive='1')
cm.var$table
```

```
##           Reference
## Prediction  0  1
##           0 67  0
##           1  0 32
```

```
# print metrics
mod1.CMmetrics <- c(cm.var$overall[c(1)], cm.var$byClass[c(1,2,5,6,7)])


mod1.CMmetrics
```

```
##      Accuracy Sensitivity Specificity Precision Recall      F1
##           1           1           1           1           1           1
```

```
# ROC and AUC
par(pty="s")
roc.stepwise <- roc(train$gentoo, model1$fitted.values, plot=TRUE, print.auc=TRUE)
```

```
## Setting levels: control = 0, case = 1
```

```
## Setting direction: controls < cases
```



```
# Dispersion Statistic
E2 <- resid(model1, type = "pearson")
N <- nrow(train)
p <- length(coef(model1)) + 1 # '+1' is due to theta
mod1.dispersion <- dispersion <- sum(E2^2) / (N - p)
```

# Multinomial Logistic Regression

```
# Initial walk-through: https://stats.idre.ucla.edu/r/dae/multinomial-logistic-regression/
```

```
# Start with initial dataset
```

```
mlr_data <- penguins
```

```
summary(mlr_data)
```

```
##      species      island bill_length_mm bill_depth_mm
## Adelie   :152  Biscoe   :168   Min.    :32.10   Min.    :13.10
## Chinstrap: 68  Dream    :124   1st Qu.:39.23   1st Qu.:15.60
## Gentoo   :124  Torgersen: 52   Median :44.45   Median :17.30
##
##                               Mean    :43.92   Mean    :17.15
##                               3rd Qu.:48.50   3rd Qu.:18.70
##                               Max.    :59.60   Max.    :21.50
##                               NA's    :2       NA's    :2
## flipper_length_mm body_mass_g      sex      year
## Min.    :172.0     Min.    :2700   female:165   Min.    :2007
## 1st Qu.:190.0     1st Qu.:3550   male  :168   1st Qu.:2007
## Median :197.0     Median :4050   NA's   : 11   Median :2008
## Mean    :200.9     Mean    :4202                   Mean    :2008
## 3rd Qu.:213.0     3rd Qu.:4750                   3rd Qu.:2009
## Max.    :231.0     Max.    :6300                   Max.    :2009
## NA's    :2        NA's    :2
```

```
mlr_data$species2 <- relevel(mlr_data$species, ref = "Gentoo")
```

```
test <- multinom(species2 ~ body_mass_g + bill_length_mm + bill_depth_mm + flipper_length_mm + island, data = mlr_data)
```

```
## # weights:  24 (14 variable)
## initial value 375.725403
## iter  10 value 20.138634
## iter  20 value 1.481353
## iter  30 value 0.035234
## iter  40 value 0.000933
## iter  50 value 0.000233
## final value 0.000089
## converged
```

```
summary(test)
```

```
## Call:
## multinom(formula = species2 ~ body_mass_g + bill_length_mm +
##      bill_depth_mm + flipper_length_mm + island, data = mlr_data)
##
## Coefficients:
##      (Intercept) body_mass_g bill_length_mm bill_depth_mm
## Adelie         179.0566 -0.01156190    -11.117549    15.793011
## Chinstrap     -151.2783 -0.04191271     4.802821    -5.639087
##      flipper_length_mm islandDream islandTorgersen
## Adelie           0.4109570    22.48619         69.34019
```

```
## Chinstrap      0.6913637   117.34484      16.71717
##
## Std. Errors:
##      (Intercept) body_mass_g bill_length_mm bill_depth_mm
## Adelie      0.736784   0.4581624      49.98112      14.88575
## Chinstrap    0.732472   0.7914301      49.74986      14.71006
##      flipper_length_mm islandDream islandTorgersen
## Adelie      6.554627   0.732472   1.236702e-24
## Chinstrap    6.545027   0.732472   1.876044e-50
##
## Residual Deviance: 0.0001778793
## AIC: 28.00018
```

```
stargazer(test, type="text", out="test.htm")
```

```
##
## =====
##              Dependent variable:
##      -----
##              Adelie      Chinstrap
##              (1)         (2)
##      -----
## body_mass_g      -0.012      -0.042
##                  (0.458)      (0.791)
##
## bill_length_mm   -11.118       4.803
##                  (49.981)     (49.750)
##
## bill_depth_mm     15.793      -5.639
##                  (14.886)     (14.710)
##
## flipper_length_mm  0.411       0.691
##                  (6.555)     (6.545)
##
## islandDream       22.486***     117.345***
##                  (0.732)     (0.732)
##
## islandTorgersen   69.340***     16.717***
##                  (0.000)     (0.000)
##
## Constant          179.057***    -151.278***
##                  (0.737)     (0.732)
##
## -----
## Akaike Inf. Crit.    28.000      28.000
## =====
## Note:                *p<0.1; **p<0.05; ***p<0.01
```

```
test.rrr = exp(coef(test))
test.rrr
```

```
##      (Intercept) body_mass_g bill_length_mm bill_depth_mm
## Adelie  5.798378e+77   0.9885047   1.484942e-05  7.224660e+06
```

```
## Chinstrap 1.998367e-66 0.9589535 1.218537e+02 3.556113e-03
## flipper_length_mm islandDream islandTorgersen
## Adelie 1.508260 5.829443e+09 1.300353e+30
## Chinstrap 1.996436 9.166779e+50 1.820439e+07
```

```
stargazer(test, type="text", coef=list(test.rrr), p.auto=FALSE, out="testrrr.htm")
```

```
##
## =====
##
## -----
##                               Adelie
##                               (1)
## -----
## body_mass_g                    0.989
##                               (0.458)
##
## bill_length_mm                0.00001
##                               (49.981)
##
## bill_depth_mm                 7,224,660.000
##                               (14.886)
##
## flipper_length_mm              1.508
##                               (6.555)
##
## islandDream                    5,829,443,146.000***
##                               (0.732)
##
## islandTorgersen                1,300,352,545,036,810,191,042,951,249,920.000***
##                               (0.000)
##
## Constant                      579,837,803,916,479,821,887,587,742,166,729,353,269,199,669,971,417,058,436,431,37
##                               (0.737)
##
## -----
## Akaike Inf. Crit.              28.000
## =====
## Note:
```

```
# Again with https://www.r-bloggers.com/2020/05/multinomial-logistic-regression-with-r/
```

```
index <- createDataPartition(mlr_data$species, p = .70, list = FALSE)
train <- mlr_data[index,]
test <- mlr_data[-index,]
```

```
# Set the reference
train$species <- relevel(train$species, ref = "Adelie")
```

```
# Training the multinomial model
multinom_model <- multinom(species ~ ., data = mlr_data)
```

```
## # weights: 36 (22 variable)
```

```
## initial value 365.837892
## iter 10 value 34.181055
## iter 20 value 0.295760
## iter 30 value 0.004662
## final value 0.000042
## converged
```

```
# Checking the model
summary(multinom_model)
```

```
## Call:
## multinom(formula = species ~ ., data = mlr_data)
##
## Coefficients:
##      (Intercept) islandDream islandTorgersen bill_length_mm bill_depth_mm
## Chinstrap -0.01000283 -0.5092332      4.737343      -4.717127      -3.9325176
## Gentoo    0.07204990 -20.1161533      7.279638      -2.880746      0.2195242
##      flipper_length_mm body_mass_g  sexmale      year species2Adelie
## Chinstrap      -1.515657  0.03478204 6.250623 0.2131091      -64.27652
## Gentoo         -1.860254  0.03139549 6.219207 0.2080230      -120.84300
##      species2Chinstrap
## Chinstrap      93.909336
## Gentoo         -8.168343
##
## Std. Errors:
##      (Intercept) islandDream islandTorgersen bill_length_mm bill_depth_mm
## Chinstrap  0.00415839 10.228172539      10.22439      82.08249      4.939571
## Gentoo    0.01810751  0.005181329      11.95099      237.58065      47.894200
##      flipper_length_mm body_mass_g  sexmale      year species2Adelie
## Chinstrap      54.22441  0.5054492 0.00447937 3.557639      2.745676
## Gentoo        204.40986  4.5030675 0.18580467 12.667205      11.926995
##      species2Chinstrap
## Chinstrap      2.749341047
## Gentoo         0.001456202
##
## Residual Deviance: 8.466626e-05
## AIC: 44.00008
```

```
# Convert the coefficients to odds by taking the exponential of the coefficients.
exp(coef(multinom_model))
```

```
##      (Intercept) islandDream islandTorgersen bill_length_mm bill_depth_mm
## Chinstrap  0.990047 6.009562e-01      114.1306  0.008940826  0.01959428
## Gentoo    1.074709 1.835125e-09      1450.4627  0.056092884  1.24548403
##      flipper_length_mm body_mass_g  sexmale      year species2Adelie
## Chinstrap  0.2196639  1.035394 518.3357 1.237520  1.216360e-28
## Gentoo    0.1556331  1.031894 502.3045 1.231241  3.300276e-53
##      species2Chinstrap
## Chinstrap  6.085640e+40
## Gentoo    2.834874e-04
```

```
head(round(fitted(multinom_model), 2))
```

```
##      Adelie Chinstrap Gentoo
## 1         1          0       0
## 2         1          0       0
## 3         1          0       0
## 5         1          0       0
## 6         1          0       0
## 7         1          0       0
```

```
# Predicting and validating the model
```

```
# Predicting the values for train dataset
```

```
train$speciesPredicted <- predict(multinom_model, newdata = train, "class")
```

```
# Building classification table
```

```
tab <- table(train$species, train$speciesPredicted)
```

```
# Calculating accuracy - sum of diagonal elements divided by total obs
```

```
round((sum(diag(tab))/sum(tab))*100,2)
```

```
## [1] 100
```

```
# Predicting the class for test dataset
```

```
test$speciesPredicted <- predict(multinom_model, newdata = test, "class")
```

```
# Building classification table
```

```
tab <- table(test$species, test$speciesPredicted)
```

```
tab
```

```
##
##      Adelie Chinstrap Gentoo
## Adelie      43          0       0
## Chinstrap    0         20       0
## Gentoo       0          0      34
```

## Thoughts

For binary, group Adelie and Chinstrap

Sex doesn't matter, Considering that I'm grouping Orange and Purple, probably don't use bill\_length\_mm, as that one shows purple and green have similar distribution

## Variables

species: of course (Y variable) ... Gentoo or Not Gentoo (No NA) island: consider it (No NA) bill\_length\_mm: don't use for binary (2 NAs) bill\_depth\_mm: use it (2 NAs) flipper\_length\_mm: use it (2 NAs) body\_mass\_g: use it (2 NAs) sex: nah (11 NAs) year: meh (No NA)

## Prompt

Data – 622 Homework # 1 Due date Feb 19, 2021- 11:59 EST Let's use the Penguin dataset for our assignment. To learn more about the dataset, please visit: <https://allisonhorst.github.io/palmerpenguins/articles/intro.html> For this assignment, let us use 'species' as our outcome or the dependent variable. 1. Logistic Regression with a binary outcome. (40) a. The penguin dataset has 'species' column. Please check how many categories you have in the species column. Conduct whatever data manipulation you need to do to be able to build a logistic regression with binary outcome. Please explain your reasoning behind your decision as you manipulate the outcome/dependent variable (species). b. Please make sure you are evaluating the independent variables appropriately in deciding which ones should be in the model. c. Provide variable interpretations in your model. 2. For your model from #1, please provide: AUC, Accuracy, TPR, FPR, TNR, FNR (20) 3. Multinomial Logistic Regression. (40) a. Please fit it a multinomial logistic regression where your outcome variable is 'species'. b. Please be sure to evaluate the independent variables appropriately to fit your best parsimonious model. c. Please be sure to interpret your variables in the model. 4. Extra credit: what would be some of the fit statistics you would want to evaluate for your model in question #3? Feel free to share whatever you can provide. (10)