

DATA 622 Assignment 2

CUNY: Spring 2021

Philip Tanofsky

15 March 2021

Introduction

```
# Import required R libraries
library(palmerpenguins)
library(tidyverse)
library(caret)
library(MASS)
library(ggplot2)
library(mvtnorm)
theme_set(theme_classic())
```

```
ds <- penguins
```

```
head(ds)
```

```
## # A tibble: 6 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g sex
##   <fct>   <fct>         <dbl>         <dbl>         <int>         <int> <fct>
## 1 Adelie  Torge~           39.1           18.7           181           3750 male
## 2 Adelie  Torge~           39.5           17.4           186           3800 fema~
## 3 Adelie  Torge~           40.3            18           195           3250 fema~
## 4 Adelie  Torge~            NA            NA            NA            NA <NA>
## 5 Adelie  Torge~           36.7           19.3           193           3450 fema~
## 6 Adelie  Torge~           39.3           20.6           190           3650 male
## # ... with 1 more variable: year <int>
```

```
summary(ds)
```

```
##      species      island  bill_length_mm  bill_depth_mm
## Adelie   :152  Biscoe   :168  Min.      :32.10  Min.      :13.10
## Chinstrap: 68  Dream    :124  1st Qu.:39.23  1st Qu.:15.60
## Gentoo   :124  Torgersen: 52  Median :44.45  Median :17.30
##                                     Mean   :43.92  Mean    :17.15
##                                     3rd Qu.:48.50  3rd Qu.:18.70
##                                     Max.    :59.60  Max.    :21.50
##                                     NA's    :2      NA's    :2
## flipper_length_mm  body_mass_g      sex      year
```

```
## Min.      :172.0      Min.      :2700      female:165      Min.      :2007
## 1st Qu.:190.0      1st Qu.:3550      male :168      1st Qu.:2007
## Median :197.0      Median :4050      NA's  : 11      Median :2008
## Mean    :200.9      Mean    :4202                      Mean    :2008
## 3rd Qu.:213.0      3rd Qu.:4750                      3rd Qu.:2009
## Max.    :231.0      Max.    :6300                      Max.    :2009
## NA's    :2          NA's    :2
```


```
dim(ds)
```

```
## [1] 344    8
```

```
glimpse(ds)
```

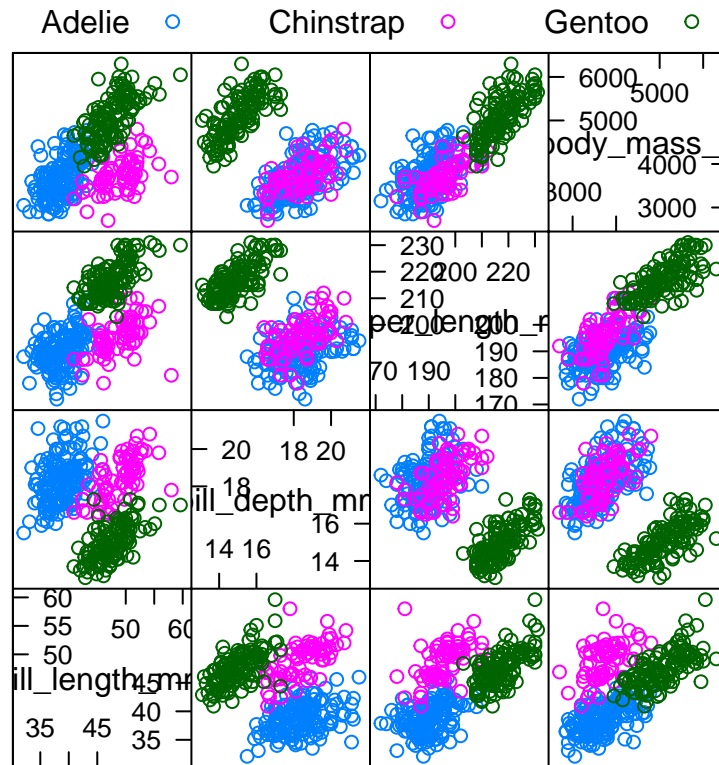
```
## Rows: 344
## Columns: 8
## $ species      <fct> Adelie, Adelie, Adelie, Adelie, Adelie, Adelie, A...
## $ island       <fct> Torgersen, Torgersen, Torgersen, Torgersen, Torge...
## $ bill_length_mm <dbl> 39.1, 39.5, 40.3, NA, 36.7, 39.3, 38.9, 39.2, 34....
## $ bill_depth_mm <dbl> 18.7, 17.4, 18.0, NA, 19.3, 20.6, 17.8, 19.6, 18....
## $ flipper_length_mm <int> 181, 186, 195, NA, 193, 190, 181, 195, 193, 190, ...
## $ body_mass_g   <int> 3750, 3800, 3250, NA, 3450, 3650, 3625, 4675, 347...
## $ sex          <fct> male, female, female, NA, female, male, female, m...
## $ year         <int> 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2007, 2...
```

```
visdat::vis_dat(ds)
```



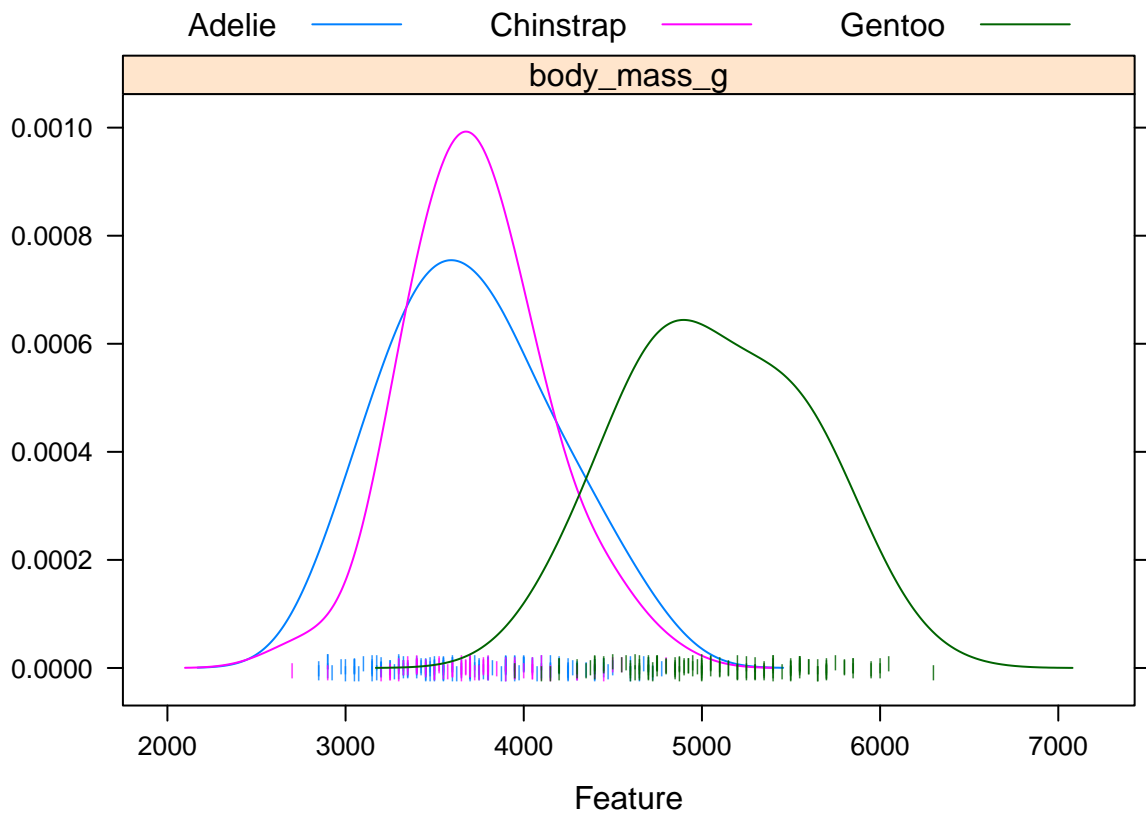
```
# Use featurePlot
# https://topepo.github.io/caret/visualizations.html

# Scatterplot
featurePlot(x = penguins[, 3:6],
            y = penguins$species,
            plot = "pairs",
            # Add a key at the top
            auto.key = list(columns = 3))
```



Scatter Plot Matrix

```
# Overlaid density plots
featurePlot(x = penguins[, 6:6],
            y = penguins$species,
            plot = "density",
            # Pass in options to xyplot() to
            # make it prettier
            scales = list(x = list(relation="free"),
                          x = list(relation="free")),
            adjust = 1.5,
            pch = "|",
            layout = c(1, 1),
            auto.key = list(columns = 3))
```




From scatterplot above featurePlot function, bill_length_mm appears to differentiate bill_depth_mm doesn't differentiate well,


```
# Plots of individual variables
#http://www.sthda.com/english/articles/32-r-graphics-essentials/133-plot-one-variable-frequency-graph-d

a <- ggplot(penguins, aes(x = bill_length_mm))


a + geom_histogram(bins = 30, color = "black", fill = "gray") +
  geom_vline(aes(xintercept = mean(bill_length_mm)),
    linetype = "dashed", size = 0.6)
```




```
b <- ggplot(penguins, aes(x = bill_depth_mm))  
  
b + geom_histogram(bins = 30, color = "black", fill = "gray") +  
  geom_vline(aes(xintercept = mean(bill_depth_mm)),  
             linetype = "dashed", size = 0.6)
```



```
c <- ggplot(penguins, aes(x = flipper_length_mm))  
c + geom_histogram(bins = 30, color = "black", fill = "gray") +  
  geom_vline(aes(xintercept = mean(flipper_length_mm)),  
             linetype = "dashed", size = 0.6)
```



```
d <- ggplot(penguins, aes(x = body_mass_g))  
  
d + geom_histogram(bins = 30, color = "black", fill = "gray") +  
  geom_vline(aes(xintercept = mean(body_mass_g)),  
             linetype = "dashed", size = 0.6)
```




Covariance checks

```
# Compute correlation matrix
cor_mat <- cor(penguins[,3:6], use = "complete.obs")
round(cor_mat, 2)
```

```
##               bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
## bill_length_mm             1.00         -0.24              0.66         0.60
## bill_depth_mm             -0.24          1.00             -0.58        -0.47
## flipper_length_mm          0.66         -0.58              1.00         0.87
## body_mass_g                0.60         -0.47              0.87         1.00
```

```
library(corrplot)
corrplot(cor_mat, type = "upper", order = "hclust",
          tl.col = "black", tl.srt = 45)
```

```
# Compute covariance matrix
cov_mat <- cov(penguins[,3:6], use = "complete.obs")
round(cov_mat, 2)
```

```
##               bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
## bill_length_mm           29.81         -2.53             50.38      2605.59
## bill_depth_mm           -2.53           3.90            -16.21     -747.37
## flipper_length_mm        50.38          -16.21           197.73     9824.42
## body_mass_g             2605.59        -747.37          9824.42    643131.08
```

LDA: Linear Discriminant Analysis

LDA does not handle categorical data well.

LDA assumes the feature variables come from multivariate normal distribution, all of them continuous

<http://www.sthda.com/english/articles/36-classification-methods-essentials/146-discriminant-analysis-essentials-in-r/>

LDA assumes the predictors are normally distributed (Gaussian distribution) and that the different classes have class-specific means and equal variance/covariance

Make sure each variable is normally distributed

<https://web.stanford.edu/class/stats202/notes/Classification/LDA.html> That is, within each class the features have multivariate normal distribution with center depending on the class and common covariance

```

# Load the data
data("penguins")

# Only complete entries
penguins <- na.omit(penguins)

# Remove 'year' and 'sex' feature
# Apparently leaving 'island' in for LDA improves the model
drops <- c("year", "sex")
penguins <- penguins[ , !(names(penguins) %in% drops)]

# Split the data into training (80%) and test set (20%)
set.seed(123)
training.samples <- penguins$species %>%
  createDataPartition(p = 0.8, list=FALSE)
train.data <- penguins[training.samples, ]
test.data <- penguins[-training.samples, ]

#2. Normalize the data. Categorical variables are automatically ignored.
# Estimate preprocessing parameters
preproc.param <- train.data %>%
  preProcess(method = c("center", "scale"))

# Transform the data using the estimated parameters
train.transformed <- preproc.param %>% predict(train.data)
test.transformed <- preproc.param %>% predict(test.data)

```

Focus on normal distributions

If n is small and the distribution of the predictors X is approximately normal in each of the classes, the LD model is more stable than logistic regression

Correlation will cause a line in the gaussian density plot

When there are K classes, linear discriminant analysis can be viewed exactly in a $K-1$ dimensional plot. Measuring which centroid is the closest. Distance in the subspace

Co-variance matrix would be 4×4 for 4 features

```

# Fit the model
model <- lda(species~., data = train.transformed)
# Make predictions
predictions <- model %>% predict(test.transformed)

# Confusion matrix
table(predictions$class, test.transformed$species)

```

```

##
##           Adelie Chinstrap Gentoo
## Adelie      29          0         0
## Chinstrap   0          13         0
## Gentoo      0          0        23

```


```
# Model accuracy
mean(predictions$class == test.transformed$species)
```

```
## [1] 1
```

```
# Output Model
model
```

```
## Call:
## lda(species ~ ., data = train.transformed)
##
## Prior probabilities of groups:
##   Adelie Chinstrap   Gentoo
## 0.4365672 0.2052239 0.3582090
##
## Group means:
##           islandDream islandTorgersen bill_length_mm bill_depth_mm
## Adelie           0.4017094           0.3247863      -0.9555507      0.6009140
## Chinstrap        1.0000000           0.0000000       0.9088301      0.6564651
## Gentoo           0.0000000           0.0000000       0.6438935     -1.1084638
##           flipper_length_mm body_mass_g
## Adelie           -0.7707877     -0.6092842
## Chinstrap        -0.3716006     -0.5901947
## Gentoo            1.1522937      1.0806975
##
## Coefficients of linear discriminants:
##               LD1          LD2
## islandDream    -1.3494973 -1.6233797
## islandTorgersen -1.1326926 -0.2217997
## bill_length_mm   0.3074563 -2.2954800
## bill_depth_mm   -1.9480393  0.3330398
## flipper_length_mm 1.1902763  0.2251057
## body_mass_g      0.9771679  0.9796534
##
## Proportion of trace:
##      LD1      LD2
## 0.8236 0.1764
```

```
# Display model
plot(model)
```



```
names(predictions)
```

```
## [1] "class"      "posterior" "x"
```

```
# Predicted classes
```

```
head(predictions$class, 6)
```

```
## [1] Adelie Adelie Adelie Adelie Adelie Adelie
```

```
## Levels: Adelie Chinstrap Gentoo
```

```
# Predicted probabilities of class membership
```

```
head(predictions$posterior, 6)
```


```
##      Adelie      Chinstrap      Gentoo
## 1 0.9999999 1.122150e-07 1.063675e-22
## 2 0.9999990 9.659679e-07 7.497505e-16
## 3 1.0000000 1.680975e-08 2.464359e-27
## 4 1.0000000 1.864923e-09 1.913351e-19
## 5 1.0000000 1.785146e-11 2.750650e-19
## 6 1.0000000 5.858899e-09 1.029719e-14
```

```
# Linear discriminants
```

```
head(predictions$x, 3)
```

```
##          LD1      LD2
## 1 -4.516164  1.880100
## 2 -2.737947  1.636759
## 3 -5.705103  2.133684
```

```
# Plot
lda.data <- cbind(train.transformed, predict(model)$x)
ggplot(lda.data, aes(LD1, LD2)) +
  geom_point(aes(color = species))
```



```
# Model accuracy
mean(predictions$class==test.transformed$species)
```

```
## [1] 1
```

```
sum(predictions$posterior[,1] >= .5)
```

```
## [1] 29
```

```
# QDA
```

```
# Remove 'island' feature as it was causing rank deficiency in group Chinstrap
#drops <- c("flipper_length_mm", "body_mass_g", "island")
drops <- c("island")
```

```
train.transformed <- train.transformed[ , !(names(train.transformed) %in% drops)]
```

```
# Fit the model
```

```
model <- qda(species~., data = train.transformed)
```

```
# Output model results
```

```
model
```

```
## Call:
```

```
## qda(species ~ ., data = train.transformed)
```

```
##
```

```
## Prior probabilities of groups:
```

```
##      Adelie Chinstrap      Gentoo
```

```
## 0.4365672 0.2052239 0.3582090
```

```
##
```

```
## Group means:
```

```
##           bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
```

```
## Adelie           -0.9555507           0.6009140          -0.7707877    -0.6092842
```

```
## Chinstrap          0.9088301           0.6564651          -0.3716006    -0.5901947
```

```
## Gentoo             0.6438935          -1.1084638           1.1522937     1.0806975
```

```
# Make predictions
```

```
predictions <- model %>% predict(test.transformed)
```

```
# Model accuracy
```

```
mean(predictions$class == test.transformed$species)
```

```
## [1] 0.9538462
```

<https://www.geeksforgeeks.org/linear-discriminant-analysis-in-r-programming/>

```
# this isn't being applied to the real data
```

```
# Variance Covariance matrix for random bivariate gaussian sample
```

```
var_covar <- matrix(data = c(1.5, 0.4, 0.4, 1.5), nrow=2)
```

```
# Random bivariate Gaussian samples for class +1
```

```
Xplus1 <- rmvnorm(400, mean = c(5, 5), sigma = var_covar)
```

```
# Random bivariate Gaussian samples for class -1
```

```
Xminus1 <- rmvnorm(600, mean = c(3, 3), sigma = var_covar)
```

```
# Samples for the dependent variable
```

```
Y_samples <- c(rep(1, 400), rep(-1, 600))
```

```
# Combining the independent and dependent variables into a dataframe
```


```
dataset <- as.data.frame(cbind(rbind(Xplus1, Xminus1), Y_samples))
```

```
colnames(dataset) <- c("X1", "X2", "Y")
```

```
dataset$Y <- as.character(dataset$Y)
```

```
# Plot the above samples and color by class labels
```

```
ggplot(data = dataset) + geom_point(aes(X1, X2, color = Y))
```



QDA: Quadratic Discriminant Analysis

Same link as above

QDA: works well with fewer features, that's when NB works well, works with higher number of features
mixed features can be used for NB

NB: Naive Bayes

<https://www.r-bloggers.com/2018/01/understanding-naive-bayes-classifier-using-r/>

```
library(e1071)

# Next load the Titanic dataset
data("Titanic")

# Save into a data frame and view it
t_df <- as.data.frame(Titanic)

# Creating data from table
repeating_sequence <- rep.int(seq_len(nrow(t_df)), t_df$Freq)
```

```

# Create the dataset by row repetition created
t_ds <- t_df[repeating_sequence, ]

# We no longer need the frequency, drop the feature
t_ds$Freq = NULL

# Fitting the Naive Bayes model
nbm <- naiveBayes(Survived~., data=t_ds)
# Output the model
nbm

# Prediction on the dataset
nb_predictions <- predict(nbm, t_ds)
# Confusion matrix to check accuracy
table(nb_predictions, t_ds$Survived)

# Getting started with Naive Bayes in mlr
library(mlr)

# Create a classification task for learning on Titanic Dataset and specify the target feature
task <- makeClassifTask(data = t_ds, target="Survived")

# Initialize the Naive Bayes classifier
selected_model <- makeLearner("classif.naiveBayes")

# Train the model
nb_mlr <- train(selected_model, task)

# Read the model learned
nb_mlr$learner.model

# Predict on the dataset without passing the target feature
predictions_mlr <- as.data.frame(predict(nb_mlr, newdata = t_ds[,1:3]))

# Confusion matrix to check accuracy
table(predictions_mlr[,1], t_ds$Survived)

```

<https://www.geeksforgeeks.org/naive-bayes-classifier-in-r-programming/>

==== Prompt =====

Homework # 2 (Generative Models) (100 points) Due on March 12, 11:59pm EST

We will be working with the Penguin dataset again as we did for Homework #1. Please use "Species" as your target variable. Using the target variable, Species, please conduct:

a. **LinearDiscriminantAnalysis**(30points):

- You want to evaluate all the 'features' or dependent variables and see what should be in your model. Please comment on the results.
- Just as a suggestion: You might want to consider exploring featurePlot on the caret package. Basically, you look at the relationship between the predictor variables and the target variable.
- Fit your LDA model using whatever predictor variables you deem appropriate. Feel free to split the data into training and testing sets.
- Look at the fit statistics/accuracy rates.

b. **QuadraticDiscriminantAnalysis**(30points)

- Same steps as above to consider

c. Naive Bayes (30 points)

a. Same steps as above to consider

d. Comment on the models fits/strength/weakness/accuracy for all these three models that you worked with. (10 p